



Statistics of Financial Markets



Chapter 2: Statistical Distributions for Financial Markets

Daniel Traian PELE

Bucharest University of Economic Studies

IDA Institute Digital Assets

Blockchain Research Center

AI4EFin Artificial Intelligence for Energy Finance

Romanian Academy, Institute for Economic Forecasting

MSCA Digital Finance

Learning Objectives

By the end of this chapter, you will be able to:

1. **Explain** why the choice of distribution affects financial risk estimation
2. **Compute** log-returns and interpret their statistical properties
3. **Identify** stylized facts of returns (heavy tails, skewness, clustering)
4. **Compare** Normal, Student- t , stable, and EVT distributions
5. **Apply** Value-at-Risk and Expected Shortfall under different distributional assumptions
6. **Evaluate** GARCH models with non-Gaussian innovations

Contents

1. Introduction
2. From Prices to Log>Returns
3. The Normal Distribution
4. Stylized Facts
5. The Student- t Distribution
6. Asymmetric Distributions
7. Stable Distributions
8. Extreme Value Theory
9. Distribution Comparison
10. Risk Management
11. Modern Extensions

Why Does the Distribution Choice Matter?

- If we assume the wrong distribution, we **underestimate** the probability of extreme losses
- **Example:** October 19, 1987 — the S&P 500 fell 20.5% in a single day
 - ▶ Under the Normal assumption: a $\sim 25\sigma$ event with probability $\approx 10^{-135}$
 - ▶ Yet it happened — the Normal distribution is inadequate for tail events
- Banks and regulators need to quantify “how bad can it get?”
- This requires **accurate distributional models**

Key Insight

- The choice of probability distribution directly determines how much capital a bank must hold against potential losses

The Distribution Choice Problem

- **Preview:** In Section 2.10 we will formally define **Value-at-Risk (VaR)**, a key risk measure that depends entirely on the distributional assumption
- The candidates we will study:
 - ▶ **Normal** — the classical benchmark (thin tails)
 - ▶ **Student- t** — heavier tails, finite higher moments
 - ▶ **Stable** — heaviest tails, infinite variance possible
 - ▶ **EVT** — models only the tails, agnostic about the center
- Each model implies a different answer to “how bad can it get?”

Model	Tail Type	Kurtosis	Parameters
Normal	Exponential	3	2
Student- t	Power law	> 3	3
Stable	Power law	∞	4
GPD (EVT)	Flexible	—	2

Chapter Roadmap

Foundations

1. From prices to log-returns
2. Normal distribution (benchmark)
3. Stylized facts of financial returns
4. Student- t distribution
5. Asymmetric distributions

Advanced Models & Applications

6. Stable distributions
7. Extreme Value Theory
8. Distribution comparison
9. **Risk management (VaR, ES)**
10. Modern extensions (GARCH)


Distribution Choice and Derivative Pricing

- ▣ Black–Scholes assumes log-normality $\Rightarrow$ constant implied volatility
- ▣ Empirical deviations produce the **volatility smile/skew**
- ▣ Portfolio optimization under non-Gaussian returns: mean-variance is suboptimal
- ▣ Regulatory context: Basel II $\rightarrow$ Basel III/IV framework
 - ▶ Basel II: 1% VaR with Normal assumption
 - ▶ Basel III/IV: shift from 1% VaR to 2.5% ES

Remark 1

The entire Basel regulatory capital framework rests on distributional assumptions. Underestimating tail risk $\Rightarrow$ insufficient capital $\Rightarrow$ systemic fragility.

S&P 500: What the Real World Looks Like



- Crashes ($|r_t| > 3\sigma$) occur far more frequently than the Normal distribution predicts [SFM_ch2_stylized_facts](#)
- Extreme returns **cluster**; we need models with heavy tails

Normal Distribution vs. Reality — Visualization and Example

(ext. Ch. 1)

Black Monday — October 19, 1987

S&P 500 fell by 20.5% in a single day.


- Under the Normal assumption ($\sigma \approx 1\%$): a $\sim 20\sigma$ event with $P \approx 10^{-89}$
- Under Student- t ($\nu=4$): $P \approx 10^{-5} — 10^{84} \times$ more plausible

Return Period

- $T = \frac{1}{252 \times 10^{-89}} \approx 4 \times 10^{86}$ years; age of the universe: $\approx 1.4 \times 10^{10}$ years
- $\Rightarrow \sim 10^{76}$ **ages of the universe!**
- The distribution choice is not a technical detail — it is a **business decision**

 SFM_ch2_normal_distribution

Normal Distribution vs. Reality — Chart



SFM_ch2_normal_distribution



Model Uncertainty and Robust Risk Measures

- ▣ **Model uncertainty:** Which distribution is “correct”? Perhaps none.
- ▣ **Ambiguity aversion** (Ellsberg paradox in portfolio choice)
- ▣ Robust risk measures under distributional uncertainty:

$$\rho_{\text{robust}}(X) = \sup_{Q \in \mathcal{Q}} \rho_Q(X)$$

where $\mathcal{Q}$ is an ambiguity set of plausible distributions

- ▣ Minimax approaches: worst-case distributional analysis
- ▣ **Model risk quantification:** range of VaR/ES across plausible models provides a “model risk premium”

Definition of Log-Returns

Definition 1 (Log-Return)

The **log-return** (continuously compounded return) at time t is:

$$r_t = \ln\left(\frac{P_t}{P_{t-1}}\right) = \ln P_t - \ln P_{t-1}$$

where P_t is the asset price at time t .

- ▣ **Additivity:** Multi-period returns are sums: $r_{t:t+k} = \sum_{i=1}^k r_{t+i}$
- ▣ **Approximation:** For small returns, $r_t \approx (P_t - P_{t-1})/P_{t-1}$
- ▣ **Statistical motivation:** Additive structure enables standard probability theory (CLT, stable limits)

Properties of Log-Returns

Why log-returns?

- Time-additive: $r_{[0,T]} = \sum_{t=1}^T r_t$
- Symmetric: +5% and -5% are comparable
- Unbounded below: can model large losses
- Enable CLT-based inference

Simple vs. log-returns

- Simple return: $R_t = \frac{P_t - P_{t-1}}{P_{t-1}}$
- Relationship: $r_t = \ln(1 + R_t)$
- Simple returns are *not* additive across time
- But simple returns *are* additive across portfolio assets

Convention

- Throughout this chapter, r_t denotes the **log-return** unless stated otherwise

Continuous-Time Limit

- Geometric Brownian Motion: $dS/S = \mu dt + \sigma dW$
- By Itô's lemma: $d \ln S = (\mu - \sigma^2/2) dt + \sigma dW$
- Jensen's inequality: $\mathbb{E}[\ln(P_T/P_0)] \neq \ln(\mathbb{E}[P_T/P_0])$
 - ▶ The arithmetic mean return exceeds the geometric mean return
 - ▶ Difference $\approx \sigma^2/2$ ("volatility drag")
- Log-return vs. simple return: implications for portfolio aggregation
 - ▶ Log-returns: additive across *time*
 - ▶ Simple returns: additive across *assets*
 - ▶ Neither is additive in both dimensions simultaneously

Volatility Drag — Demonstration

Thought experiment

- ▣ Your portfolio gains 50% today and loses 50% tomorrow
- ▣ Arithmetic mean return: $(+50\% - 50\%)/2 = 0\%$. But the real money?
- ▣ $\$100 \rightarrow \$150 \rightarrow \$75$. You lost **25%**!

Why log-returns tell the truth

- ▣ $r_1 = \ln(1.5) = +40.5\%$, $r_2 = \ln(0.5) = -69.3\%$
- ▣ Sum: $r_1 + r_2 = -28.8\% = \ln(0.75)$ ✓ (additivity works!)
- ▣ Simple return is **not** additive: $(+50\%) + (-50\%) \neq \ln(0.75)$
- ▣ **Rule:** $\mathbb{E}[\text{geometric return}] \approx \mathbb{E}[\text{arithmetic return}] - \sigma^2/2$; high volatility $\Rightarrow$ more severe real loss

Semimartingale Representation

- Price processes modeled as **semimartingales**: $S_t = S_0 + A_t + M_t$
- Itô's lemma for semimartingales with jumps:

$$d \ln S_t = \frac{dS_t}{S_{t-}} - \frac{1}{2} \frac{d[S, S]_t^c}{S_{t-}^2} + \sum_{s \leq t} \left(\ln \frac{S_s}{S_{s-}} - \frac{\Delta S_s}{S_{s-}} \right)$$

- Quadratic variation and realized volatility:

$$[X, X]_t = \lim_{\|\Pi\| \rightarrow 0} \sum_i (X_{t_{i+1}} - X_{t_i})^2$$

- Non-continuous price processes: jumps and Lévy–Itô decomposition
- High-frequency returns reveal microstructure noise, bid-ask bounce, Epps effect

“The Prince of Mathematicians”



Carl Friedrich Gauss
(1777–1855)

From stars to Wall Street

- ▣ Derived the normal curve in 1809 for **astronomical errors**
- ▣ The idea: errors are the sum of many small, independent perturbations
- ▣ Appears on the **10 DM banknote** — the only distribution on a banknote!

The warning

- ▣ Gauss derived it for measurement errors, **not** for financial markets
- ▣ Wall Street adopted it anyway — the consequences will be seen in this chapter

The Only Distribution on a Banknote



- 10 Deutsche Mark (series IV, 1993–2001): Gauss + normal curve with formula $f(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$
- On the reverse: Gauss's sextant and the geodetic triangulation network of Germany

The Normal Distribution — Definition and Properties

Definition 2 (Normal Distribution)

$X \sim N(\mu, \sigma^2)$ if:

$$f(x) = \frac{1}{\sigma\sqrt{2\pi}} \exp\left(-\frac{(x-\mu)^2}{2\sigma^2}\right), \quad x \in \mathbb{R}.$$


Moments:

- ▣ $\mathbb{E}[X] = \mu, \quad \text{Var}(X) = \sigma^2$
- ▣ Skewness = 0
- ▣ Kurtosis = 3 (excess kurtosis = 0)

Key properties:

- ▣ Symmetric about μ
- ▣ 68–95–99.7 rule
- ▣ Tails decay as $e^{-x^2/2}$

The Normal Distribution — PDF Chart

 SFM_ch2_normal_distribution

- **Gauss's bell**: symmetric, unimodal shape, maximum at $x = \mu$
- Parameters μ and σ^2 control the **location** and **spread**
- Tails decay **exponentially**: $f(x) \sim e^{-x^2/2}$ — underestimates extreme events

Quantiles and the Central Limit Theorem

- ▣ **Quantiles:** The α -quantile of $N(\mu, \sigma^2)$ is:

$$q_\alpha = \mu + \sigma \Phi^{-1}(\alpha)$$

This will be used in Section 2.10 to define Value-at-Risk

- ▣ **Key quantiles:** $\Phi^{-1}(0.01) = -2.326$, $\Phi^{-1}(0.05) = -1.645$


Theorem 1 (Central Limit Theorem)

If $X_1, X_2, \dots$ are i.i.d. with $\mathbb{E}[X_i] = \mu$ and $\text{Var}(X_i) = \sigma^2 < \infty$, then:

$$\frac{\bar{X}_n - \mu}{\sigma/\sqrt{n}} \xrightarrow{d} N(0, 1).$$

- ▣ **Stability:** Sum of independent normals is normal: $X + Y \sim N(\mu_1 + \mu_2, \sigma_1^2 + \sigma_2^2)$


CLT Illustration



- As n increases, $\bar{X}_n$ approaches normality
- Limitation:** CLT requires finite variance — what if $\sigma^2 = \infty$?

Normal Fit to S&P 500 Returns

SFM_ch2_normal_test



- Normal PDF fits the center well but **misses the tails**
- Histogram has more mass in the tails than Normal predicts

Maximum Entropy and Normality Tests

- Maximum entropy: Among all distributions with given (μ, σ^2) , the Normal maximizes Shannon entropy


$$H = - \int f \ln f \, dx$$

- Normality tests:

Jarque–Bera test:

$$JB = \frac{n}{6} \left(S^2 + \frac{(K-3)^2}{4} \right) \sim \chi^2(2)$$

Tests skewness S and excess kurtosis $K - 3$ jointly.



Other tests:

- Shapiro–Wilk (regression-based)
- Anderson–Darling (CDF, tail-weighted)
- Kolmogorov–Smirnov (sup CDF distance)

B. Normality Test Results

Test	Statistic	p-value (Critical)	Reject H ₀ (5%)
Jarque-Bera	38879.71	p = 0.00e+00	Yes
Shapiro-Wilk	0.90985	p = 1.64e-47	Yes
Anderson-Darling	126.296	adj(p) = 0.7020	Yes

Sample size: n = 0.031
 Skewness: 0.3491 | Excess kurtosis: 10.8342
 H₀: Data follows a Normal distribution

The Jarque–Bera Test

- ▣ **Maximum entropy:** Among all distributions with given (μ, σ^2) , the Normal maximizes Shannon entropy
 $H = -\int f \ln f \, dx$

Hypotheses

- H_0 : $S = 0$ and $K = 3$ (the data come from a Normal distribution)
 H_1 : $S \neq 0$ or $K \neq 3$ (departure from normality)

Test statistic

$$JB = \frac{n}{6} \left(S^2 + \frac{(K - 3)^2}{4} \right) \xrightarrow{d} \chi^2(2)$$

- ▣ $S = \frac{1}{n} \sum \left(\frac{x_i - \bar{x}}{s} \right)^3$ (skewness), $K = \frac{1}{n} \sum \left(\frac{x_i - \bar{x}}{s} \right)^4$ (kurtosis)
- ▣ Reject H_0 if $JB > \chi^2_{2, 1-\alpha}$. For $\alpha = 5\%$: $\chi^2_{2, 0.95} = 5.99$

The Shapiro–Wilk Test

Shapiro–Wilk Test (1965)

H_0 : the data come from a Normal distribution vs. H_1 : non-normality

$$W = \frac{(\sum_{i=1}^n a_i x_{(i)})^2}{\sum_{i=1}^n (x_i - \bar{x})^2}, \quad W \in (0, 1]$$

- ▣ $x_{(i)}$ = order statistics of the sample
- ▣ a_i = tabulated coefficients, depending on n
- ▣ Reject H_0 if W is too small (close to 0)
- ▣ The most powerful test for small samples ($n < 50$)

The Anderson–Darling and Kolmogorov–Smirnov Tests

Anderson–Darling Test (1954)

$$H_0: F = F_0 \quad \text{vs.} \quad H_1: F \neq F_0$$

$$A^2 = -n - \frac{1}{n} \sum_{i=1}^n (2i-1) [\ln F_0(x_{(i)}) + \ln(1 - F_0(x_{(n+1-i)}))]$$

- Tail-weighted: more sensitive to deviations in the extremes of the distribution than KS


Kolmogorov–Smirnov Test (1933)

$$\square \quad H_0: F = F_0 \quad \text{vs.} \quad H_1: F \neq F_0$$

$$D_n = \sup_x |F_n(x) - F_0(x)|$$

- Non-parametric; less sensitive to tails than AD

Normality Tests — Visualization



B. Normality Test Results

Test	Statistic	p-value / Critical	Reject H ₀ (5%)
Jarque-Bera	30576.71	p = 0.00e+00	Yes
Shapiro-Wilk	0.909685	p = 1.64e-47	Yes
Anderson-Darling	126.2996	cv(5%) = 0.7520	Yes

Sample size: n = 6,537
 Skewness: -0.3489 | Excess kurtosis: 10.6242
 H₀: Data follows a Normal distribution

 SFM_ch2_normal_distribution

- The test results on S&P 500 returns **massively** reject the normality hypothesis
- JB test: $p\text{-value} \approx 0$ (skewness and kurtosis differ from Normal values)

Normality Tests — Empirical Results

- Moment generating function: $M_X(t) = \exp(\mu t + \sigma^2 t^2/2)$
- Multivariate normal: $\mathbf{X} \sim N(\boldsymbol{\mu}, \boldsymbol{\Sigma})$ with Mahalanobis distance

Test	Statistic	p -value	Decision
Jarque–Bera	> 1000	< 0.001	Reject H_0
Shapiro–Wilk	< 0.99	< 0.001	Reject H_0
Anderson–Darling	> 10	< 0.001	Reject H_0

Conclusion

- Normality is overwhelmingly rejected for daily financial returns
- The Normal distribution underestimates tail risk

Berry–Esseen Bounds and Convergence Rates

Theorem 2 (Berry–Esseen)

If $X_1, X_2, \dots$ are i.i.d. with $\mathbb{E}[X_i] = 0$, $\mathbb{E}[X_i^2] = \sigma^2$, $\mathbb{E}[|X_i|^3] = \rho < \infty$, then:

$$\sup_x |F_n(x) - \Phi(x)| \leq \frac{C \rho}{\sigma^3 \sqrt{n}}$$

where $C \leq 0.4748$ (Shevtsova, 2011).

- ▣ Rate of convergence: $O(n^{-1/2})$ — slow for heavy-tailed data
- ▣ Edgeworth expansions refine the CLT approximation for finite n
- ▣ **Normal as stable:** $N(\mu, \sigma^2) = S(2, 0, \sigma/\sqrt{2}, \mu)$ — the $\alpha = 2$ member of the stable family
- ▣ When $\mathbb{E}[X_i^2] = \infty$: CLT fails, but the **Generalized CLT** gives stable limits

The Pre-Asymptotic Regime — Why the CLT Does Not Save Us

The Berry–Esseen $\leftrightarrow$ aggregational Gaussianity connection

- ▣ The CLT **promises** convergence toward normality as we aggregate
- ▣ But Berry–Esseen shows the speed depends on ρ/σ^3 (third absolute moment)
- ▣ Financial returns have **very large** ρ/σ^3 (heavy tails) $\Rightarrow$ convergence is **painfully slow**

Practical implication

- ▣ If returns had thin tails, 22 days (one month) would suffice for perfect normality
- ▣ Due to heavy tails, even **annual** aggregates are not perfectly Gaussian
- ▣ **Conclusion** (Taleb): in finance, we permanently live in the *pre-asymptotic regime* — the CLT is a promise our data cannot honor

The Man Who Broke the Bell Curve



Benoît Mandelbrot
(1924–2010)

From cotton to fractals

- **1963**: cotton prices had variations *far* too large for the Normal
- Proposed **Lévy stable** distributions ($\alpha < 2$) — heavy tails, infinite variance
- Inventor of **fractal geometry**; self-similarity at every scale
- “*Markets are riskier than people think*”

The legacy

- Ignored for 40 years by Wall Street, which preferred the Gaussian model
- The 2008 crisis proved him right — thin-tailed models failed spectacularly

Fat Tails and Excess Kurtosis

- ▣ **Kurtosis:** $\kappa = \mathbb{E}[(X - \mu)^4]/\sigma^4$; excess kurtosis: $\kappa_e = \kappa - 3$
- ▣ Normal distribution: $\kappa_e = 0$
- ▣ Typical daily equity kurtosis: $\kappa \in [5, 50]$, so $\kappa_e \gg 0$
- ▣ **Interpretation:** Extreme returns occur far more often than the Normal predicts

Asset	Skewness	Excess Kurtosis
S&P 500 (daily)	-0.3	≈ 10
FTSE 100 (daily)	-0.2	≈ 8
USD/EUR (daily)	-0.1	≈ 5
Bitcoin (daily)	-0.5	≈ 15

Negative Skewness and Volatility Clustering

Negative skewness

- ▣ Left tail heavier than right tail
- ▣ Especially pronounced for equity indices
- ▣ Crashes are sharper than rallies
- ▣ Skewness $S = \mathbb{E}[(X - \mu)^3]/\sigma^3 < 0$

Volatility clustering

- ▣ “Large changes followed by large changes” (Mandelbrot, 1963)
- ▣ $\text{Corr}(|r_t|, |r_{t+k}|) > 0$ for large k
- ▣ Returns are uncorrelated but **not independent**
- ▣ $\text{Corr}(r_t, r_{t+k}) \approx 0$ but $\text{Corr}(r_t^2, r_{t+k}^2) > 0$

Implication

- ▣ A single i.i.d. distribution cannot capture volatility clustering
- ▣ We need time-varying models (GARCH, Section 2.11)

Aggregational Gaussianity

- ▣ **Key observation:** As the holding period increases, the distribution of returns becomes *more Gaussian*
- ▣ Daily returns: high kurtosis, heavy tails
- ▣ Weekly returns: moderate kurtosis
- ▣ Monthly returns: closer to Normal
- ▣ **Mechanism:** CLT-like effect — monthly return = $\sum$ of ~ 22 daily returns

Frequency	Excess Kurtosis	JB p -value
Daily	≈ 10	< 0.001
Weekly	≈ 4	< 0.001
Monthly	≈ 1.5	≈ 0.05
Quarterly	≈ 0.5	≈ 0.3


Summary of Stylized Facts

1. **Fat tails:** $\kappa_e \gg 0$ — extreme events more frequent than Normal
2. **Negative skewness:** Left tail heavier (equities)
3. **Volatility clustering:** Periods of high/low volatility persist
4. **Aggregational Gaussianity:** Kurtosis decreases with horizon
5. **No return autocorrelation:** $\text{Corr}(r_t, r_{t+k}) \approx 0$ for $k \geq 1$
6. **Squared return autocorrelation:** $\text{Corr}(r_t^2, r_{t+k}^2) > 0$

The Challenge

- Any adequate model for financial returns must reproduce these stylized facts
- The Normal distribution fails on points 1, 2, and 3

Visualizing Volatility Clustering



SFM_ch2_stylized_facts

- “Large changes followed by large changes” (Mandelbrot, 1963) — the 2001, 2008, 2020 crises
- $\text{Corr}(r_t, r_{t+k}) \approx 0$ but $\text{Corr}(|r_t|, |r_{t+k}|) > 0$ — uncorrelated, but **not independent**

Stylized Facts Equations — Kurtosis and Autocorrelation

Kurtosis

$$\kappa = \frac{\mathbb{E}[(X - \mu)^4]}{\sigma^4}$$

- ▣ **Excess kurtosis:** $\kappa_e = \kappa - 3$
- ▣ $\kappa_e = 0$: Normal; $\kappa_e > 0$: **leptokurtic**

Autocorrelation of r_t^2

$$\rho_k = \text{Corr}(r_t^2, r_{t+k}^2) > 0, \quad k \text{ large}$$

Numerical example

	κ	κ_e
Normal	3	0
S&P 500	≈ 10	≈ 7
FTSE 100	≈ 8	≈ 5
Bitcoin	≈ 15	≈ 12

- ▣ $\kappa_e = 7$ for S&P 500 $\Rightarrow$ tails **substantially heavier** than Normal

Preview: Risk Measures — VaR and ES

Definition 3 (Value-at-Risk — working definition)

Value-at-Risk at level α : the maximum loss that is **not** exceeded with probability $(1 - \alpha)$:

$$\text{VaR}_\alpha = -F^{-1}(\alpha), \quad \text{where } F^{-1} \text{ is the return quantile.}$$


Definition 4 (Expected Shortfall — working definition)

Expected Shortfall (ES) at level α : the **average** loss in the worst $\alpha\%$ scenarios:

$$\text{ES}_\alpha = -\frac{1}{\alpha} \int_0^\alpha F^{-1}(p) dp.$$

- ▣ **Example:** Normal 1% VaR, $\sigma = 1.5\%$, \$1M: $\text{VaR} = 2.326 \times 0.015 \times 10^6 = \$34,890$
- ▣ Formal properties, Artzner coherence, backtesting — **Section 10**

VaR and ES — Graphical Illustration



SFM_ch2_risk_management

- ▣ VaR_α : quantile — “we will not lose more than X with prob. $1 - \alpha$ ”
- ▣ ES_α : tail mean — “if we exceed VaR, how much do we lose *on average*?”
- ▣ $\text{ES} \geq \text{VaR}$ always; the gap grows with tail heaviness

Leverage Effect and Long Memory

Leverage effect

- ▣ Negative returns increase future volatility *more* than positive returns of equal magnitude
- ▣ Asymmetric response: $\text{Corr}(r_t, \sigma_{t+1}^2) < 0$
- ▣ Named after Black (1976): falling stock price $\Rightarrow$ higher leverage $\Rightarrow$ higher risk
- ▣ Captured by GJR-GARCH, EGARCH

Long memory in volatility

- ▣ $\text{Corr}(|r_t|, |r_{t+k}|)$ decays *slowly* (hyperbolically, not exponentially)
- ▣ Hurst exponent $H > 0.5$ for $|r_t|$
- ▣ Standard GARCH: exponential decay
- ▣ FIGARCH, HAR models capture long memory

Gain/Loss Asymmetry

- ▣ **Gain/loss asymmetry:** Stock prices decline faster than they rise
- ▣ Large drawdowns happen in days; recovery takes months or years
- ▣ This asymmetry is distinct from skewness — it concerns the *dynamics*, not the marginal distribution
- ▣ **Absence of return autocorrelation:** $\text{Corr}(r_t, r_{t+k}) \approx 0$ for $k \geq 1$
 - ▶ Markets are approximately efficient at the mean level
 - ▶ But dependence exists in higher moments (r_t^2 , $|r_t|$)

Cont (2001) Taxonomy and Multifractal Models

- Cont (2001): systematic taxonomy of 11 stylized facts for asset returns
- **Multifractal models** (Mandelbrot, Fisher & Calvet, 1997):
 - ▶ $\mathbb{E}[|r_t|^q] \propto \Delta t^{\zeta(q)}$ with *nonlinear* $\zeta(q)$ (multiscaling)
 - ▶ Unifractal: $\zeta(q) = qH$ (linear); multifractal: $\zeta(q)$ concave
- Cross-correlation structure and multivariate stylized facts
- Intraday stylized facts: microstructure noise, bid-ask bounce, Epps effect

Fractal Structure and Power Laws

- ▣ **Power-law tails:** $P(|r| > x) \sim x^{-\alpha}$ with $\alpha \in [2, 5]$ for most financial assets
- ▣ Tail exponent α is remarkably stable across assets and time periods
- ▣ **Self-similarity:** Return distributions at different time scales are related by scaling
- ▣ Connection to stable distributions: if $\alpha < 2$, the GCLT implies stable limits
- ▣ The “inverse cubic law” (Gopikrishnan et al., 1999): $\alpha \approx 3$ for US equities
- ▣ Debate: true power law vs. slowly varying tails

Who Is “Student”?



William Sealy Gosset
(1876–1937)

The story behind the pseudonym

- Chemist at the **Guinness** brewery, Dublin
- The problem: how do you test beer quality from **very small** samples? ($n = 3-4$)
- Guinness **forbade** publication — trade secret!
- The solution: publish under the pseudonym “**Student**” (1908)

Why does it matter for finance?

- The Normal with σ *estimated* is no longer Normal — it is Student- t !
- The irony: invented for beer, now it models stock market crashes

Student- t Distribution — DefinitionDefinition 5 (Student- t Distribution)


$X \sim t_\nu$ if:

$$f(x; \nu) = \frac{\Gamma\left(\frac{\nu+1}{2}\right)}{\sqrt{\nu\pi} \Gamma\left(\frac{\nu}{2}\right)} \left(1 + \frac{x^2}{\nu}\right)^{-(\nu+1)/2}, \quad x \in \mathbb{R}.$$


- ▣ Single parameter $\nu > 0$ (**degrees of freedom**) controls tail heaviness
- ▣ $\mathbb{E}[X] = 0$ ($\nu > 1$); $\text{Var}(X) = \frac{\nu}{\nu-2}$ ($\nu > 2$)
- ▣ Excess kurtosis: $\kappa_e = \frac{6}{\nu-4}$ ($\nu > 4$)
- ▣ $\nu \rightarrow \infty$: converges to $N(0, 1)$; $\nu = 1$: Cauchy distribution
- ▣ Tails: $f(x) \sim |x|^{-(\nu+1)}$ (**power law**)

Student- t PDFs — Effect of ν

SFM_ch2_student_t_fit



- Smaller $\nu \Rightarrow$ heavier tails, higher peak
- Financial returns: typically $\hat{\nu} \in [3, 8]$ for daily equity data

Student- t Fit to S&P 500

- Student- t captures the heavy tails much better than the Normal
- Estimated $\hat{\nu} \approx 4-5$ for daily S&P 500; center and tails well-matched

Why Student- t for Finance?

Property	Normal	Student- t
Tail decay	Exponential ($e^{-x^2/2}$)	Power law ($ x ^{-\nu-1}$)
Excess kurtosis	0	$6/(\nu - 4)$
Parameters	2 (μ, σ)	3 (μ, σ, ν)
Extreme events	Very rare	More frequent
Limit	—	Normal ($\nu \rightarrow \infty$)

Key Advantage


- Student- t **neests** the Normal as a special case ($\nu \rightarrow \infty$) while allowing heavier tails
- It is the simplest heavy-tailed alternative

Location-Scale Form and MLE

- ▣ **Location-scale form:** $X = \mu + \sigma T$, where $T \sim t_\nu$
- ▣ MLE estimation of (μ, σ, ν) : maximize


$$\ell(\mu, \sigma, \nu) = \sum_{t=1}^n \ln f\left(\frac{r_t - \mu}{\sigma}; \nu\right) - n \ln \sigma$$

- ▣ Information criteria for model selection:
 - ▶ AIC = $-2\ell + 2k$ (penalizes complexity lightly)
 - ▶ BIC = $-2\ell + k \ln n$ (penalizes complexity more)
- ▣ Student- t as innovation distribution in GARCH models: $\epsilon_t = \sigma_t z_t$, $z_t \sim t_\nu$
- ▣ QQ-plot diagnostics: departures from the 45° line indicate misspecification

Model Comparison: Normal vs. Student- t 

Model	Log-lik	AIC	BIC
Normal (μ, σ)	ℓ_1	$-2\ell_1 + 4$	$-2\ell_1 + 2 \ln n$
Student- t (μ, σ, ν)	$\ell_2 > \ell_1$	$-2\ell_2 + 6$	$-2\ell_2 + 3 \ln n$

- Student- t achieves substantially better AIC/BIC despite one extra parameter

Normal vs. Student- t — Overlay on S&P 500[SFM_ch2_distribution_comparison](#)

- Student- t captures the heavy tails; the Normal systematically underestimates them

Model Comparison: AIC/BIC

Model	Log-lik.	AIC	BIC
Normal (μ, σ)	ℓ_1	$-2\ell_1 + 4$	$-2\ell_1 + 2 \ln n$
Student- t (μ, σ, ν)	$\ell_2 > \ell_1$	$-2\ell_2 + 6$	$-2\ell_2 + 3 \ln n$

- Student- t achieves substantially better AIC/BIC despite one extra parameter
- The log-likelihood gain comes from the superior tail fit

Mixture Representation and Bayesian Interpretation

- **Mixture representation:** $X | V \sim N(0, V)$, $V \sim \text{Inv-Gamma}(\nu/2, \nu/2)$
 - ▶ Student- t = Normal with random variance (scale mixture of normals)
 - ▶ Interpretation: V captures time-varying volatility
- **Bayesian interpretation:** Student- t as marginal posterior under conjugate normal–inverse-gamma prior
- **Multivariate Student- t :** $f(\mathbf{x}) \propto \left(1 + \frac{(\mathbf{x}-\boldsymbol{\mu})' \boldsymbol{\Sigma}^{-1} (\mathbf{x}-\boldsymbol{\mu})}{\nu}\right)^{-(\nu+p)/2}$
- **Tail comparison:** t_ν tail index = ν vs. stable tail index = α
- Student- t is **not** closed under convolution; stable distributions **are**

Model Uncertainty and Robust Risk Measures

- ▣ **Model uncertainty:** Which distribution is “correct”? Perhaps none — every model is an *approximation*.
- ▣ **Ambiguity aversion** (Ellsberg paradox): the decision-maker does not know the exact distribution and adopts a conservative attitude.
- ▣ Robust risk measures under distributional uncertainty:

$$\rho_{\text{robust}}(X) = \sup_{Q \in \mathcal{Q}} \rho_Q(X)$$

- ▣ **Ambiguity set** $\mathcal{Q}$ = a set of probability distributions deemed plausible. We compute risk under the *worst-case* distribution from $\mathcal{Q}$.
- ▣ Classic examples of $\mathcal{Q}$:
 - ▶ **Moment-based:** all distributions with given (μ, σ^2) : $\mathcal{Q} = \{Q : \mathbb{E}_Q[X] = \mu, \text{Var}_Q(X) = \sigma^2\}$
 - ▶ **Wasserstein:** neighborhood of radius ε around the empirical distribution: $\mathcal{Q} = \{Q : W_p(Q, \hat{P}_n) \leq \varepsilon\}$
 - ▶ **KL-divergence:** $\mathcal{Q} = \{Q : D_{\text{KL}}(Q \| P_0) \leq \eta\}$
- ▣ **Model risk quantification:** the VaR/ES range across plausible models provides a “model risk premium”

Distances Between Distributions — Intuition

- **Problem:** How do we measure how “close” two distributions P and Q are?
- **Wasserstein distance** $W_p(P, Q)$:
 - ▶ Also known as *Earth Mover's Distance*
 - ▶ Intuition: the minimum cost of “transporting” probability mass from P to Q
 - ▶ If P and Q are distributions on $\mathbb{R}$ with CDFs F_P, F_Q :

$$W_1(P, Q) = \int_{-\infty}^{+\infty} |F_P(x) - F_Q(x)| dx$$

- ▶ Advantage: accounts for the *geometry* of the space (not just density overlap)
- **Kullback–Leibler divergence** $D_{\text{KL}}(Q \| P_0)$:
 - ▶ Measures the “information loss” when approximating Q by P_0
 - ▶ $D_{\text{KL}}(Q \| P_0) = \int q(x) \ln \frac{q(x)}{p_0(x)} dx \geq 0$, with equality iff $Q = P_0$
 - ▶ **Not** a proper distance (not symmetric: $D_{\text{KL}}(Q \| P) \neq D_{\text{KL}}(P \| Q)$)
- **Application:** The choice of radius ε (Wasserstein) or η (KL) controls the *degree of conservatism* of the robust risk measure — larger $\varepsilon \Rightarrow$ wider $\mathcal{Q} \Rightarrow$ higher robust VaR

Wasserstein in Practice — Bank Stress-Testing

Scenario: The 2008 Crisis / 2020 Pandemic

- ▣ The *historical* return distribution suddenly becomes irrelevant (non-stationarity)
- ▣ VaR estimated on 2005–2007 data massively underestimates losses in 2008
- ▣ **Problem:** how do we protect ourselves when the model we estimated may be wrong?

Solution: robust optimization with Wasserstein neighborhood

- ▣ Instead of computing VaR under a *single* distribution $\hat{P}$, compute under the **worst-case** distribution in a neighborhood: $\rho_{\text{robust}} = \sup_{Q: w_1(Q, \hat{P}) \leq \varepsilon} \text{VaR}_Q$
- ▣ The radius ε = the bank's uncertainty budget (how much reality can deviate)
- ▣ Small ε : optimism (“the crisis will not repeat”); large ε : conservatism (survival)
- ▣ **Basel III+**: regulators mandate *stressed VaR* — exactly this philosophy

From Student- t to Pareto — Why Do We Need Both?

Student- t : the **complete** model

- ▣ Models the **entire dataset** (center + tails)
- ▣ Trade-off: fitting the center well may miss the extreme tail
- ▣ Global parameter: ν simultaneously controls the peak and tails


Pareto: the **extreme tail** model

- ▣ Models **only the extreme values** ($x > x_m$)
- ▣ No compromises: focuses exclusively on the right tail
- ▣ Mentally prepares for Extreme Value Theory (EVT)

The mathematical link

- ▣ Student- t with ν d.f. has **Pareto tails**: $P(X > x) \sim C \cdot x^{-\nu}$ for $x \rightarrow \infty$
- ▣ The Pareto tail index α = the degrees of freedom ν of Student- t
- ▣ Pareto is not an alternative to Student- t , but a **complementary** diagnostic tool

The 80/20 Rule — The Man Behind the Law



Vilfredo Pareto
(1848–1923)

From gardening to inequality

- ▣ Italian-Swiss engineer: **80% of Italy's land** was owned by **20% of the population** (1896)
- ▣ Legend: 20% of plants produced 80% of peas
- ▣ The same law appears **everywhere**: wealth, cities, earthquakes, sales, web traffic

Why does it matter for finance?

- ▣ A few catastrophic days generate most of the risk
- ▣ Mandelbrot (1963): return tails are Pareto — $P(X > x) \propto x^{-\alpha}$

The Pareto Distribution — Power Law

Definition (Pareto, 1896)

A random variable X follows a **Pareto distribution** with parameters $x_m > 0$ (scale) and $\alpha > 0$ (shape / tail index) if:

$$f(x) = \frac{\alpha x_m^\alpha}{x^{\alpha+1}}, \quad x \geq x_m$$

▣ **CDF:** $F(x) = 1 - \left(\frac{x_m}{x}\right)^\alpha$

▣ **Survival function:**

$$\bar{F}(x) = P(X > x) = \left(\frac{x_m}{x}\right)^\alpha$$

▣ **Power-law** behavior: $P(X > x) \propto x^{-\alpha}$


▣ **Mean:** $\mathbb{E}[X] = \frac{\alpha x_m}{\alpha-1}$ (exists only if $\alpha > 1$)

▣ **Variance:** $\text{Var}(X) = \frac{\alpha x_m^2}{(\alpha-1)^2(\alpha-2)}$ (exists only if $\alpha > 2$)

▣ $\alpha \leq 2$: **infinite variance**

▣ $\alpha \leq 1$: **infinite mean**

The Pareto Distribution — PDF and CDF

 SFM_ch2_pareto_analysis

- Small $\alpha \Rightarrow$ heavy tail, density decays slowly
- Large $\alpha \Rightarrow$ mass concentrates near x_m , tail decays rapidly

The Pareto Distribution — Power-Law Signature

 SFM_ch2_pareto_analysis


Log-log graphical test

- On a log-log plot, the Pareto survival function is a **straight line**:

$$\log P(X > x) = -\alpha \log x + \alpha \log x_m$$

- **Slope** = $-\alpha$ (tail index)
- Alignment on a line $\Rightarrow$ Pareto behavior; deviation for small x : *truncation*
- Practical check: plot $\log \hat{S}(x)$ vs. $\log x$ on empirical data — linearity confirms a Pareto tail

The Pareto Distribution — Log-log Survival



The Pareto Distribution — Why It Matters in Finance

The Pareto principle (80/20 rule)

- ▣ **Vilfredo Pareto** (1896): 80% of Italy's wealth was held by 20% of the population
- ▣ The same law appears in finance: extreme losses follow a power law
- ▣ **Mandelbrot** (1963): market returns have Pareto-type tails

Key property: heavy tails


- ▣ Normal: $P(X > x) \sim e^{-x^2/2}$ (**exponential**); Pareto: $P(X > x) \sim x^{-\alpha}$ (**polynomial**)
- ▣ *Black Swan* events are **much more probable** under Pareto
- ▣ **Example**: $\alpha = 3$, $P(X > 10x_m) = 10^{-3}$; under Normal: $\approx 10^{-23}$

The Pareto Distribution — Parameter Estimation (MLE)

Maximum Likelihood Estimator (MLE)

- ▣ Let $x_1, \dots, x_n$ be an i.i.d. sample from $\text{Pareto}(\alpha, x_m)$
- ▣ **Step 1:** $\hat{x}_m = \min(x_1, \dots, x_n)$ (scale = minimum value)
- ▣ **Step 2:** $\hat{\alpha}_{\text{MLE}} = \frac{n}{\sum_{i=1}^n \ln(x_i / \hat{x}_m)}$
- ▣ **Properties:** consistent, asymptotically efficient, $\sqrt{n}$ -convergent



The Pareto Distribution — MLE: Demonstration



□ **Left:** histogram + MLE fit vs. true value ($\alpha = 3$, $n = 1000$)

□ **Right:** MLE convergence as n increases ($n = 50 \rightarrow 5000$)


The Pareto Distribution — Hill Estimator on S&P 500



SFM_ch2_pareto_analysis

- **Hill plot:** stable region $\hat{\alpha} \approx 2.9$; **Tail fit:** Pareto $\hat{\alpha} \approx 2.5$
- $\alpha < 3 \Rightarrow$ moments of order ≥ 3 do not exist; confirmation: **heavy tails**

The Pareto Distribution — Tail Fitting (S&P 500)



 SFM_ch2_pareto_analysis

- **QQ-plot:** alignment on diagonal $\Rightarrow$ tail is Pareto
- **Threshold sensitivity:** stable region $\approx 3\text{--}8\%$; $\hat{\alpha} \approx 2.5\text{--}3.0$

The Pareto Distribution — Tail Fitting Interpretation

Tail fitting procedure

- ▣ **Step 1:** Select threshold u (or tail fraction, e.g. top 5%)
- ▣ **Step 2:** Estimate $\hat{\alpha}$ on data $\{x_i : x_i > u\}$ via MLE or Hill
- ▣ **Step 3:** Check sensitivity: vary u and observe $\hat{\alpha}$ stability
- ▣ **Step 4:** Validate with QQ-plot and CDF fit

Empirical results (S&P 500, 2000–2025)

- ▣ $\hat{\alpha} \approx 2.5\text{--}3.0 \Rightarrow$ finite moments only up to order 2
- ▣ Variance exists, but kurtosis **does not** — the Normal distribution is inadequate
- ▣ The Pareto CDF fits well on the right tail of losses

The Pareto Distribution — Applications

 SFM_ch2_pareto_analysis



Zipf's law (market capitalization)

- ▣ Firm capitalization follows a **power law**: $P(\text{cap} > x) \propto x^{-\alpha}$
- ▣ Top 10 firms $\approx$ 30% of the market; Zipf slope ≈ -1 (universal empirical law)

Other applications of the Pareto distribution

- ▣ **Insurance losses**: extreme insured amounts $\Rightarrow$ Pareto
- ▣ **Operational risk** (Basel II): operational losses have Pareto tails
- ▣ **Wealth distribution**: Top 1% holds 40–50% of wealth (Piketty, 2014)
- ▣ **Trading volume**: a few large trades dominate total volume
- ▣ **Cyber risk**: severity of cyber attacks is Pareto-distributed

Zipf's Law — Market Capitalization



The Pareto Distribution — Theoretical Connections

Connection with EVT and the Hill estimator

- ▣ **Pickands–Balkema–de Haan theorem:** exceedances above a sufficiently high threshold approximately follow a **Generalized Pareto Distribution** (GPD)
- ▣ The classical Pareto distribution = special case of GPD with $\xi > 0$
- ▣ **Hill estimator:** estimates α from order statistics $X_{(1)} \leq \dots \leq X_{(n)}$:

$$\hat{\alpha}_{\text{Hill}} = \left(\frac{1}{k} \sum_{i=1}^k \ln \frac{X_{(n-i+1)}}{X_{(n-k)}} \right)^{-1}$$

Connection with stable distributions and Student- t

- ▣ α -stable distributions with $\alpha < 2$ have Pareto tails: $P(X > x) \sim C_{\alpha} x^{-\alpha}$
- ▣ Student- t with ν degrees of freedom: Pareto tails with index $\alpha = \nu$
- ▣ The Pareto tail index unifies the theory: smaller $\alpha \Rightarrow$ heavier tails

Case Study: The Data


SFM_ch2_case_study_2a

Prices and log-returns

- ▣ **S&P 500** — developed market, high liquidity
- ▣ **BRD** — Romanian stock, emerging market
- ▣ **Bitcoin** — digital asset, extreme volatility
- ▣ Period: 2015–2025, daily data

Asset	N	Mean (%)	Std (%)	Skew	Kurt
S&P 500	~2 500	0.05	1.15	−0.60	13.5
BRD	~2 500	0.04	1.10	−0.80	12.0
Bitcoin	~3 600	0.12	3.80	−0.30	8.5

Case Study: Prices and Log-Returns



Case Study: Stylized Facts

 SFM_ch2_case_study_2a


Volatility clustering

- ▣ ACF of $|r_t|$ confirms long-range dependence
- ▣ All 3 assets exhibit significant autocorrelation of $|r_t|$ up to lag 60
- ▣ **Kurtosis** $\gg 3$ across all markets $\Rightarrow$ heavy tails
- ▣ **Negative skewness** for S&P 500 and BRD (*leverage effect*)
- ▣ Bitcoin: lower skewness, but $3\times$ higher volatility

Remark 2

- ▣ Volatility clustering invalidates the i.i.d. assumption required by the Black–Scholes model

Case Study: Stylized Facts — Chart



SFM_ch2_case_study_2a

Case Study: Normal vs Student- t

 SFM_ch2_case_study_2a


Log-scale histogram with overlay

- We compare the two models on a logarithmic scale
- Normal systematically underestimates the tails
- Student- t adapts well with estimated ν :
 - ▶ S&P 500: $\hat{\nu} \approx 4-5$; BRD: $\hat{\nu} \approx 3-5$; Bitcoin: $\hat{\nu} \approx 3-4$
- The difference grows exponentially in the tails

Implication


- On a logarithmic scale, the Normal underestimation becomes visible: empirical tails are orders of magnitude heavier

Case Study: Normal vs Student- t — Chart



SFM_ch2_case_study_2a

Case Study: QQ-Plots



SFM_ch2_case_study_2a

- Normal QQ: systematic deviations in tails Student- t QQ: much better alignment
- Bitcoin: largest deviations, confirming heavier tails

Case Study: Normality Tests

Testing H_0 : returns $\sim \mathcal{N}(\mu, \sigma^2)$

Asset	Jarque–Bera	Shapiro–Wilk	Anderson–Darling	Kolmogorov–Smirnov
S&P 500	> 10 000***	< 0.99***	> 20***	> 0.05***
BRD	> 8 000***	< 0.99***	> 15***	> 0.04***
Bitcoin	> 5 000***	< 0.98***	> 30***	> 0.06***

□ *** rejection at level $\alpha = 0.001$ for all tests and all assets

Conclusion

- Normality is **decisively rejected** across all three markets
- All four tests converge: JB (moments), SW (order), AD (tails), KS (full distribution)

Case Study: VaR and ES — Normal vs Student- t Applying risk formulas (portfolio 1M EUR, $\mu \approx 0$)

$$\square \text{ VaR}_{1\%}^N = 2.326 \sigma; \quad \text{VaR}_{1\%}^t = \sqrt{\frac{\nu-2}{\nu}} |t_{\nu}^{-1}(0.01)| \sigma; \quad \text{ES: left-tail mean below VaR}$$

Asset ($\sigma, \hat{\nu}$)	VaR Normal	VaR Student- t	ES Normal	ES Student- t
S&P 500 (1.15%, 4.5)	26 750	30 400	30 500	40 200
BRD (1.10%, 3.8)	25 590	29 600	29 300	40 200
Bitcoin (3.80%, 3.2)	88 390	102 000	100 900	148 700

Normal ES underestimates by 30–50%

$$\square P(r < -3\sigma): \text{Normal} = \Phi(-3) = 0.135\% \text{ (1 day/741); } \text{Student-}t \text{ } (\hat{\nu}=4.5) \approx 0.58\% \text{ (1 day/172)} \text{ — } 4.3\times \text{ more frequent}$$

Case Study: Model Selection — AIC/BIC

Information criteria ($AIC = -2\ell + 2k$; $BIC = -2\ell + k \ln n$)

Asset	ΔAIC (Normal – Student- t)	ΔBIC	Winner
S&P 500	> 200	> 190	Student- t
BRD	> 150	> 140	Student- t
Bitcoin	> 300	> 290	Student- t

Remark 3

- ▣ $\Delta AIC > 10 \Rightarrow$ strong evidence against the Normal model (Burnham & Anderson)
- ▣ Student- t wins decisively across all markets, despite having one extra parameter (ν)

Case Study: Conclusions

Lessons from real data analysis

1. **Kurtosis** $\gg 3$ across all markets $\Rightarrow$ the Normal model is inadequate
2. **Student- t** provides a first approximation, but has *symmetric tails*
3. **Skewness** (especially for S&P 500 and BRD) requires asymmetric models
4. **Volatility clustering** demands conditional models (GARCH)

Risk management implications

- Normal ES at 1% underestimates the mean tail loss by **30–50%** compared to Student- t
- Normal VaR underestimates by $\sim 15\%$, but 3σ events occur $4\times$ more often

AI Exercise: Distributions and Tail Risk

Prompt to test in ChatGPT / Claude / Copilot

“Download daily data for S&P 500 and Bitcoin from 2015 to 2025 using Python (yfinance). Compute log-returns. Fit a Normal and a Student- t distribution via MLE on each series. Compare models using AIC/BIC, QQ-plot, and log-scale histogram. Compute 1% VaR and Expected Shortfall under both models. Estimate the Pareto tail index using the Hill estimator.”

Exercise:

1. Does the AI compute log-returns correctly? Check: $r_t = \ln(P_t/P_{t-1})$, not $(P_t - P_{t-1})/P_{t-1}$
2. Are the estimated degrees of freedom $\hat{\nu}$ plausible ($\hat{\nu} \in [3, 8]$ for daily data)?
3. Are the VaR and ES formulas correct? Check: $\text{VaR}_{1\%}^t$ uses adjusted scale $s = \sigma\sqrt{(\nu - 2)/\nu}$?
4. Ask the AI: “What happens to VaR if $\nu \leq 2$?” Does it mention infinite variance?
5. Does the Hill plot show a stable region? What value of $\hat{\alpha}$ does it obtain? Compare with $\hat{\nu}$ from Student- t

⚠ **Warning:** AI often confuses *excess kurtosis* ($\kappa - 3$) with total kurtosis (κ). Always verify the definition used.

Question 1

Question

- ☐ The event of October 19, 1987 (S&P 500 -20.5%) corresponds to a $\sim 20\sigma$ event under the Normal assumption ($\sigma \approx 1\%$). What is the probability of such an event?

Answer choices

- ☐ (A) $\approx 10^{-2}$ (1 in 100)
- ☐ (B) $\approx 10^{-30}$ (practically impossible on a human timescale)
- ☐ (C) $\approx 10^{-89}$ (impossible in the lifetime of the universe)
- ☐ (D) $\approx 10^{-6}$ (once every million days)

Question 1: Answer

Answer: (C)

- ▣ **Normal:** $\sigma \approx 1\%$, so $-20.5\% = 20.5\sigma$. $P = 2\Phi(-20.5) \approx 10^{-89}$.
- ▣ **Student- t ($\nu = 5$):** adjusted scale $s = \sigma\sqrt{3/5} = 0.775\%$; $t_{\text{stat}} = 20.5/0.775 = 26.5$; $P(|T_5| > 26.5) \approx 10^{-8}$ — still extreme, but $10^{81} \times$ **more probable** than under Normal!
- ▣ Yet it happened — demonstrating the **inadequacy of the Normal distribution**.
- ▣ Under stable distributions ($\alpha \approx 1.7$), the probability rises to $\sim 10^{-3}$.

Question 2

Question

- ☐ A Student- t distribution with $\nu = 4$ degrees of freedom has:

Answer choices

- ☐ (A) Finite mean, finite variance, finite kurtosis
- ☐ (B) Finite mean, finite variance, infinite kurtosis
- ☐ (C) Finite mean, infinite variance, infinite kurtosis
- ☐ (D) Infinite mean, infinite variance, infinite kurtosis

Question 2: Answer

Answer: (B)

- ▣ For Student- t : mean exists if $\nu > 1$, variance if $\nu > 2$, kurtosis if $\nu > 4$.
- ▣ With $\nu = 4$: finite mean ($\checkmark$), finite variance ($= \nu/(\nu - 2) = 2$, $\checkmark$), but kurtosis **does not exist** (4th moment diverges: $6/(\nu - 4)$ with $\nu = 4$).
- ▣ (A) is wrong: the 4th moment does not exist at $\nu = 4$ exactly.
- ▣ This shows that even distributions with finite variance can have **very heavy** tails.

Question 3: Reasoning — Basel III and Distribution Choice

Question

- ▣ A bank computes regulatory capital using ES at 2.5% under the Normal distribution. Supervisors discover that portfolio returns have $\hat{\nu} = 4$ (Student- t). What is the consequence?

Answer choices

- ▣ **(A)** Capital is sufficient — ES does not depend on the distribution shape
- ▣ **(B)** Capital is underestimated by 30–50%, the bank must increase reserves
- ▣ **(C)** Capital is overestimated — Student- t predicts lower risks
- ▣ **(D)** Cannot decide without knowing the volatility

Question 3: Answer

Answer: (B)

- ▣ Student- t with $\nu = 4$ has much heavier tails than Normal: $P(r < -3\sigma)$ is $\sim 4\times$ larger.
- ▣ ES under Student- t is **30–50% larger** than under Normal (as demonstrated in the case study).
- ▣ (A) is wrong: ES depends critically on the shape of the tails, not just on σ .
- ▣ (D) is wrong: even with the same σ , Student- t gives larger ES due to heavier tails.
- ▣ **Lesson:** Distribution choice is not academic — it **directly determines bank capital**.

Exercise 1: Log>Returns and Moments

Problem

□ Given daily closing prices:

t	0	1	2	3	4
P_t	100	103	98	95	102

1. Compute log-returns $r_t = \ln(P_t/P_{t-1})$ for $t = 1, \dots, 4$.
2. Compute sample mean, variance, skewness, and kurtosis.
3. Apply the Jarque–Bera test: $JB = \frac{n}{6} \left(S^2 + \frac{(K-3)^2}{4} \right)$. Is normality rejected at 5%?
4. Compute the total log-return over 4 periods. Verify time-additivity.

Exercise 2: The Normal Distribution and VaR

Problem

□ A \$5M portfolio has daily returns $\sim N(\mu, \sigma^2)$ with $\hat{\mu} = 0.04\%$ and $\hat{\sigma} = 1.2\%$

1. Compute daily VaR at levels 5%, 1%, and 0.1%.
2. Compute Expected Shortfall at the 1% level.
3. A -4% event occurred. What is its probability under the Normal assumption?
4. Argue why Normal VaR may underestimate real risk.

Hint

$\text{VaR}_\alpha = -(\mu + z_\alpha \cdot \sigma)$, where $z_{0.05} = -1.645$, $z_{0.01} = -2.326$, $z_{0.001} = -3.090$.

Exercise 3: Student- t — Fitting and Comparison

Problem

▣ Using the data from Exercise 1, suppose MLE produces: Normal $\ell = 12.5$, Student- t ($\nu = 5$) $\ell = 14.8$

1. Compute AIC and BIC for both models ($k_N = 2$, $k_t = 3$, $n = 4$).
2. Which model is preferred by AIC? By BIC?
3. Compute the excess kurtosis of the Student- t distribution with $\nu = 5$.
4. Explain why Student- t captures heavy tails better than the Normal distribution.

Formulas

▣ $\text{AIC} = -2\ell + 2k$, $\text{BIC} = -2\ell + k \ln n$, $\kappa = 6/(\nu - 4)$ for $\nu > 4$

Exercise 4: Comparing VaR/ES — Normal, Student- t , Pareto

Problem

□ A \$10M portfolio has daily returns with $\hat{\sigma} = 1.5\%$, $\mu \approx 0$. The Hill estimator on the left tail gives $\hat{\alpha} = 3.5$

1. Compute 1% VaR under: (a) Normal; (b) Student- t with $\nu = 5$; (c) Pareto with $\alpha = 3.5$, $x_m = 0.01$
2. Compute 1% ES under each model. *Hint:* $ES_{\alpha}^{\text{Pareto}} = \frac{\alpha}{\alpha-1} \text{VaR}_{\alpha}$ for $\alpha > 1$
3. Fill in the table and discuss the differences:

	Normal	Student- t ($\nu=5$)	Pareto ($\alpha=3.5$)
VaR 1%	?	?	?
ES 1%	?	?	?
ES/VaR	?	?	?

Python Lab — Interactive Exercises

Lab 1: Distribution Fitting (Y3)

1. Open `SFM_ch2_normal_distribution.ipynb`
2. Download S&P 500 data (2015–2025); compute log-returns
3. Histogram + Normal overlay; Jarque–Bera test
4. Fit Student- t (MLE). Compare AIC with Normal model

 `SFM_ch2_normal_distribution`

Exercise 1: Solution

Solution

1. Log-returns: $r_1 = \ln(103/100) = 0.0296$, $r_2 = \ln(98/103) = -0.0498$, $r_3 = \ln(95/98) = -0.0311$, $r_4 = \ln(102/95) = 0.0713$
2. $\bar{r} = (0.0296 - 0.0498 - 0.0311 + 0.0713)/4 = 0.0050$; $s^2 = \frac{1}{3} \sum (r_i - \bar{r})^2 = 0.00253$; $s = 0.0503$.
Skewness: $S \approx 0.56$; Kurtosis: $K \approx 1.72$
3. $JB = \frac{4}{6}(0.56^2 + (1.72 - 3)^2/4) = 0.48$. With $\chi^2_2(0.05) = 5.99$: $0.48 < 5.99 \Rightarrow$ **do not reject** normality (sample too small!)
4. $r_{0:4} = \ln(102/100) = 0.0198 = r_1 + r_2 + r_3 + r_4 = 0.0198 \checkmark$

Exercise 2: Solution

Solution

1. $\text{VaR}_\alpha = -(\mu + z_\alpha \sigma) \times W$, with $\mu = 0.0004$, $\sigma = 0.012$, $W = \$5\text{M}$:
▶ $\text{VaR } 5\%: 0.01934 \times 5\text{M} = \$96,680$; $\text{VaR } 1\%: \$137,560$; $\text{VaR } 0.1\%: \$183,400$
2. $\text{ES}_{1\%} = \sigma \cdot \phi(-2.326)/0.01 = 0.03198 \Rightarrow \$159,900$
3. $P(r < -0.04) = \Phi(-3.37) \approx 0.038\%$ (1 day in $\sim 2,600$)
4. Empirically, events $> 3\sigma$ occur much more frequently $\Rightarrow$ Normal VaR underestimates risk

Exercise 3: Solution

Solution

1. With $\ell_N = 12.5$, $\ell_t = 14.8$, $k_N = 2$, $k_t = 3$, $n = 4$:
► Normal: $\text{AIC} = -21.0$; $\text{BIC} = -22.23$. Student- t : $\text{AIC} = -23.6$; $\text{BIC} = -25.44$
2. Student- t wins on both criteria ($\text{AIC}: -23.6 < -21.0$; $\text{BIC}: -25.44 < -22.23$)
3. Excess kurtosis: $\kappa = 6/(\nu - 4) = 6/(5 - 4) = 6$. Total kurtosis: $3 + 6 = 9$
4. The Student- t PDF decays as $|x|^{-(\nu+1)}$ (power law) vs. $e^{-x^2/2}$ for Normal. The slower decay allocates more probability to the tails

References I

-  Mandelbrot, B. (1963). The variation of certain speculative prices. *Journal of Business*, 36(4), 394–419.
-  Fama, E. (1965). The behavior of stock-market prices. *Journal of Business*, 38(1), 34–105.
-  Cont, R. (2001). Empirical properties of asset returns: stylized facts and statistical issues. *Quantitative Finance*, 1(2), 223–236.
-  Hansen, B. (1994). Autoregressive conditional density estimation. *International Economic Review*, 35(3), 705–730.
-  Nolan, J. P. (1997). Numerical calculation of stable densities and distribution functions. *Stochastic Models*, 13(4), 759–774.
-  McNeil, A., Frey, R. & Embrechts, P. (2015). *Quantitative Risk Management*. Princeton University Press.
-  Barndorff-Nielsen, O. (1977). Exponentially decreasing distributions for the logarithm of particle size. *Proceedings of the Royal Society of London A*, 353, 401–419.

References II

-  Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series. *Econometrica*, 57(2), 357–384.
-  Merton, R. (1976). Option pricing when underlying stock returns are discontinuous. *Journal of Financial Economics*, 3, 125–144.
-  Rachev, S. T. (2003). *Handbook of Heavy Tailed Distributions in Finance*. Elsevier.
-  Nolan, J. P. (2020). *Stable Distributions — Models for Heavy Tailed Data*. Birkhäuser.
-  Vuong, Q. H. (1989). Likelihood ratio tests for model selection and non-nested hypotheses. *Econometrica*, 57(2), 307–333.
-  Fissler, T. & Ziegel, J. F. (2016). Higher order elicibility and Osband's principle. *Annals of Statistics*, 44(4), 1680–1707.
-  Barndorff-Nielsen, O. & Shephard, N. (2001). Non-Gaussian Ornstein–Uhlenbeck-based models. *Journal of the Royal Statistical Society B*, 63(2), 167–241.

References III



Vuong, Q. H. (1989). Likelihood ratio tests for model selection and non-nested hypotheses. *Econometrica*, 57(2), 307–333.



Fissler, T. & Ziegel, J. F. (2016). Higher order elicibility and Osband's principle. *Annals of Statistics*, 44(4), 1680–1707.

Online resources and code

- ▣ **Quantlet:** <https://quantlet.com> > Code platform for quantitative methods
- ▣ **Quantinar:** <https://quantinar.com> > Learning platform for quantitative methods
- ▣ **GitHub SFM:** https://github.com/QuantLet/Stat_fin_markets/tree/master/Ch_02 > Python code for Chapter 2

References IV



Kelly, B., Jiang, H. & Pruitt, S. (2022). Tail risk and asset prices. *Review of Financial Studies*, 35(5), 2254–2295.



Bollerslev, T. & Todorov, V. (2023). The jump tails of asset returns. *Econometrica* (forthcoming).



Aas, K., Czado, C., Frigessi, A. & Bakken, H. (2009). Pair-copula constructions of multiple dependence. *Insurance: Mathematics and Economics*, 44(2), 182–198.

Recent relevant events

- ▣ **Covid-19 (Mar. 2020):** S&P 500 –12% in one day; the 3rd worst event in history
- ▣ **FTX/crypto (Nov. 2022):** –65% collapse in 48h; extreme heavy tails
- ▣ **SVB (Mar. 2023):** Banking contagion; tail risk in the financial sector
- ▣ **Lesson:** Heavy-tailed models are **essential**, not optional

Key Takeaways — What to Remember from This Chapter

5 fundamental lessons

1. The Normal distribution **systematically underestimates** tail risk
2. Student- t ($\hat{\nu} \in [3, 8]$) captures heavy tails with a single extra parameter; decisively wins on AIC/BIC
3. Pareto and power laws unify tail theory: α connects Student- t , stable distributions, and EVT
4. ES is more informative than VaR — Normal underestimates ES by 30–50%
5. **The distribution choice directly determines regulatory capital:** Basel III FRTB (2023) mandates 2.5% ES

Preparation for upcoming chapters

- ▣ **Ch. 3:** Stable distributions, skew- t , GED — asymmetry and heavier tails
- ▣ **Ch. 4:** Extreme Value Theory (EVT) — generalized Pareto, thresholds

Thank You!

Questions?

Course materials are available at: <https://danpele.github.io/SFM/>



Quantlet



Quantinar