



Statistics of Financial Markets

Seminar 3: Stable Distributions



Daniel Traian PELE
Antoaneta Amza

Bucharest University of Economic Studies
IDA Institute Digital Assets
Blockchain Research Center
AI4EFin Artificial Intelligence for Energy Finance
Romanian Academy, Institute for Economic Forecasting
MSCA Digital Finance


Seminar Contents

- ▣ **Objectives** – Learning goals and seminar roadmap
- ▣ **Part I: Manual exercises** – Stability property, characteristic function, moments and tails
- ▣ **Part II: Estimation** – McCulloch's method, MLE, Koutrouvelis regression
- ▣ **Part III: Tails and risk** – Heavy tails comparison, Gaussian vs Stable VaR
- ▣ **Part IV: Python implementation** – Simulation, fitting, VaR backtest
- ▣ **Case study: S&P 500** – Real-data stable fitting and cross-asset comparison
- ▣ **Homework** – Stable distributions on BVB
- ▣ **Summary** – Competence checklist
- ▣ **Seminar planning** – Weekly topic overview
- ▣ **Quick quiz** – 5 review questions
- ▣ **References** – Key textbooks and online resources

Learning Objectives

By the end of this seminar you will be able to:

- ▣ Verify the stability property for Gaussian and Cauchy
- ▣ Compute the characteristic function of a stable distribution
- ▣ Analyse moments and tail behaviour
- ▣ Estimate parameters of a stable distribution (McCulloch, MLE)
- ▣ Compare VaR: Gaussian vs Stable vs Empirical
- ▣ Implement stable fitting in Python



Exercise 1: Verifying Stability — The Gaussian Case

Problem

- Let X_1, X_2 be i.i.d. $\sim \mathcal{N}(0, 1)$. Show that $aX_1 + bX_2 \stackrel{d}{=} cX + d$ for $a = 3, b = 4, c = 5$.

Solution

- If $X_1 \sim \mathcal{N}(0, 1)$, then $aX_1 \sim \mathcal{N}(0, a^2)$ and $bX_2 \sim \mathcal{N}(0, b^2)$.
- Sum: $aX_1 + bX_2 \sim \mathcal{N}(0, a^2 + b^2) = \mathcal{N}(0, 9 + 16) = \mathcal{N}(0, 25)$.
- Hence $c = \sqrt{a^2 + b^2} = \sqrt{25} = 5$ and $d = 0$.
- **Conclusion:** The normal distribution is stable with $\alpha = 2$. The stability property reduces to the variance addition rule.

Exercise 1 cont.: Cauchy Case + Comparison Table

Cauchy case ($\alpha = 1$)

- ▣ Let $X_1, X_2 \sim \text{Cauchy}(0, 1)$ i.i.d.
- ▣ $aX_1 + bX_2 \sim \text{Cauchy}(0, a + b)$
- ▣ For $a = 3, b = 4$: $c = a + b = 7, d = 0$.
- ▣ **Note:** For Cauchy, scales add *linearly*, not quadratically!

General rule

- ▣ For $S(\alpha, \beta, \gamma, \delta)$:

$$c^\alpha = a^\alpha + b^\alpha$$

	α	Rule	c
Gaussian	2	$c^2 = a^2 + b^2$	5
Cauchy	1	$c = a + b$	7
$S(1.5)$	1.5	$c^{1.5} = a^{1.5} + b^{1.5}$	6.19

Exercise 2: The Characteristic Function

Problem

- Compute $\varphi_X(t)$ for $X \sim S(1.5, 0.3, 1, 0)$ at $t = 1$.

Solution

- The characteristic function of a stable distribution:

$$\log \varphi_X(t) = -|t|^\alpha \left[1 - i\beta \operatorname{sgn}(t) \tan \frac{\pi\alpha}{2} \right] + i\delta t$$

- For $\alpha = 1.5$, $\beta = 0.3$, $\gamma = 1$, $\delta = 0$, $t = 1$:

$$\begin{aligned} \log \varphi_X(1) &= -|1|^{1.5} \left[1 - i \cdot 0.3 \cdot 1 \cdot \tan \frac{3\pi}{4} \right] \\ &= -1 \cdot [1 - i \cdot 0.3 \cdot (-1)] = -1 - 0.3i \end{aligned}$$

- Therefore:

$$\varphi_X(1) = e^{-1-0.3i} = e^{-1}(\cos 0.3 - i \sin 0.3) \approx 0.351 - 0.109i$$

Exercise 3: Moments and Tails

Problem

- A financial analyst models daily returns with $S(1.7, -0.2, 0.01, 0.0005)$. Which moments exist? What are the risk management implications?

Solution

□ Moment existence:

- ▶ $\mathbb{E}[|X|^p] < \infty$ if and only if $p < \alpha$
- ▶ $\alpha = 1.7 \Rightarrow$
 - Mean ($p = 1$): **exists** ($1 < 1.7$)
 - Variance ($p = 2$): **does not exist** ($2 > 1.7$)
 - Kurtosis ($p = 4$): **does not exist**

□ Practical implications:

- ▶ $\beta = -0.2$: negative skewness (downside risk)
- ▶ Gaussian VaR **underestimates** true risk
- ▶ Probability of a -5σ event:
 - Normal: $\approx 3 \times 10^{-7}$ (1 in 3.5M days)
 - Stable(1.7): ≈ 0.003 (1 in 333 days)
- ▶ Sharpe ratio is not rigorously defined!

McCulloch's Method — Concept

Key idea

- Estimate parameters $(\alpha, \beta, \gamma, \delta)$ from **sample quantiles**.

Intermediate statistics:

$$v_{\alpha} = \frac{q_{0.95} - q_{0.05}}{q_{0.75} - q_{0.25}}$$

$$v_{\beta} = \frac{q_{0.95} + q_{0.05} - 2q_{0.50}}{q_{0.95} - q_{0.05}}$$

- v_{α} measures *tail thickness*
- v_{β} measures *asymmetry*

Reference table (selection):

v_{α}	α	Interpretation
2.44	2.0	Gaussian
2.62	1.9	Near Gaussian
3.04	1.7	Moderate tails
4.21	1.5	Heavy tails
6.31	1.3	Very heavy tails

McCulloch — Worked Example

Numerical data

- From a sample of 2,000 observations: $q_{0.05} = -3.42$, $q_{0.25} = -0.78$, $q_{0.50} = 0.05$, $q_{0.75} = 0.82$, $q_{0.95} = 3.61$.

Step-by-step solution

- Step 1: Compute v_α**

$$v_\alpha = \frac{q_{0.95} - q_{0.05}}{q_{0.75} - q_{0.25}} = \frac{3.61 - (-3.42)}{0.82 - (-0.78)} = \frac{7.03}{1.60} = 4.39$$

- Step 2: Estimate α**

► From the table: $v_\alpha = 4.39$ corresponds to $\hat{\alpha} \approx 1.48$.

- Step 3: Compute v_β**

$$v_\beta = \frac{q_{0.95} + q_{0.05} - 2q_{0.50}}{q_{0.95} - q_{0.05}} = \frac{3.61 + (-3.42) - 2 \cdot 0.05}{7.03} = \frac{0.09}{7.03} \approx 0.013$$

Maximum Likelihood Estimation (MLE)

Procedure

$$\hat{\theta}_{MLE} = \arg \max_{\theta} \sum_{i=1}^n \log f(x_i; \alpha, \beta, \gamma, \delta)$$

Challenges:

- ▶ PDF has no closed form (in general)
- ▶ Requires numerical evaluation via FFT
- ▶ Likelihood function is not convex
- ▶ Slow convergence for small α

Advantages:

- ▶ Asymptotically most efficient estimator
- ▶ Standard confidence intervals
- ▶ Modern implementation:
`scipy.stats.levy_stable.fit()`

Python one-liner:

```
from scipy.stats import levy_stable  
a, b, loc, sc = levy_stable.fit(data)
```

Koutrouvelis Regression

Two-stage method

- ▣ Uses the empirical characteristic function (ECF).

- ▣ **Stage 1:** Estimate α and γ via linear regression:

$$\log(-\log|\hat{\varphi}_n(t)|^2) = \log(2\gamma^\alpha) + \alpha \log|t|$$

- ▣ Regression $Y = a + bX$ with $Y_k = \log(-\log|\hat{\varphi}_n(t_k)|^2)$ and $X_k = \log|t_k|$.

- ▣ **Stage 2:** Estimate β and δ via regression on the imaginary part:

$$\arg[\hat{\varphi}_n(t)] = \delta t - \beta \gamma^\alpha \tan \frac{\pi\alpha}{2} \operatorname{sgn}(t)|t|^\alpha$$

Advantage

- ▣ Numerically robust, does not require PDF evaluation, fast convergence.

Heavy Tails — Gaussian vs Stable

Q SFM_ch2_stable_fit

Comparison table: Probability of extreme events

Event	Gaussian	Stable($\alpha=1.7$)	Ratio
$ X > 3\sigma$	2.7×10^{-3}	1.9×10^{-2}	$\approx 7\times$
$ X > 5\sigma$	5.7×10^{-7}	3.0×10^{-3}	$\approx 5,000\times$
$ X > 10\sigma$	1.5×10^{-23}	5.2×10^{-4}	$\approx 10^{19}\times$

Practical implication

- Gaussian models dramatically underestimate the risk of extreme events. A “ 10σ ” event that the Gaussian considers virtually impossible occurs roughly once every $\sim 2,000$ days (~ 8 years) under the stable model.

Exercise 4: Gaussian vs Stable VaR

Problem

- Portfolio of 1,000,000 EUR. Daily returns modelled with $S(1.7, -0.2, 0.01, 0.0005)$. Compare VaR at 1%, 5% and 0.1%.


Solution

VaR level	Gaussian	Stable	Difference
5%	-16,449 EUR	-21,340 EUR	+30%
1%	-22,763 EUR	-39,870 EUR	+75%
0.1%	-30,091 EUR	-87,420 EUR	+190%

- Conclusion:** The deeper into the tail, the larger the gap between Gaussian and Stable VaR. At 0.1%, the stable model estimates losses nearly **3 times larger**.

Visual Comparison: Survival Function

SFM_ch2_stable_fit



Simulating Stable Distributions

```
import numpy as np
from scipy.stats import levy_stable
import matplotlib.pyplot as plt

# Simulate stable distributions for different alpha
alphas = [0.5, 1.0, 1.5, 2.0]
x = np.linspace(-8, 8, 1000)

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(14, 5))
colors = ['#DC3545', '#FF8C00', '#2E7D32', '#1A3A6E']
labels = [r'$\alpha=0.5$ (Levy)', r'$\alpha=1.0$ (Cauchy)',
          r'$\alpha=1.5$', r'$\alpha=2.0$ (Gaussian)']

for a, c, lab in zip(alphas, colors, labels):
    pdf = levy_stable.pdf(x, alpha=a, beta=0)
    ax1.plot(x, pdf, color=c, label=lab)      # Linear scale
    ax2.semilogy(x, pdf, color=c, label=lab)  # Log scale

ax1.set_title('PDFs (linear)')
ax2.set_title('PDFs (log scale) - tail behavior')
for ax in [ax1, ax2]:
    ax.legend(loc='upper center', bbox_to_anchor=(0.5, -0.12),
              ncol=2, frameon=False)
plt.tight_layout()
```

Fitting Stable to S&P 500

```
import yfinance as yf
from scipy.stats import levy_stable, norm


# Download S&P 500 data
sp500 = yf.download('^GSPC', start='2000-01-01', end='2025-12-31')
returns = np.log(sp500['Close'] / sp500['Close'].shift(1)).dropna()

# Fit stable distribution
alpha, beta, loc, scale = levy_stable.fit(returns.values)
print(f"Stable fit: alpha={alpha:.3f}, beta={beta:.3f}")

# Fit normal
mu, sigma = norm.fit(returns.values)
print(f"Normal fit: mu={mu:.6f}, sigma={sigma:.6f}")


# Compare on log-scale
x = np.linspace(returns.min()*1.2, returns.max()*1.2, 1000)
plt.hist(returns, bins=150, density=True, alpha=0.4, label='Data')
plt.semilogy(x, levy_stable.pdf(x, alpha, beta, loc, scale),
             label=f'Stable (alpha={alpha:.2f})')
plt.semilogy(x, norm.pdf(x, mu, sigma), '--', label='Normal')
plt.legend(loc='upper center', bbox_to_anchor=(0.5, -0.1),
          ncol=3, frameon=False)
```

Overlay PDF: Histogram + Stable + Normal (log-scale)



VaR Comparison + Backtest

```
# VaR comparison
levels = [0.01, 0.05, 0.001]
for cl in levels:
    var_s = levy_stable.ppf(
        cl, alpha, beta, loc, scale)
    var_n = norm.ppf(cl, mu, sigma)
    var_e = np.percentile(
        returns, cl*100)
    print(f"VaR_{cl*100:.1f}%:")
    print(f"  Stable: {var_s:.5f}")
    print(f"  Normal: {var_n:.5f}")
    print(f"  Empirical: {var_e:.5f}")
```



AI Exercise: Prompt Engineering

Exercise

□ Use ChatGPT or Claude to generate Python code that:

1. Downloads BTC-USD data for the last 5 years
2. Fits a stable distribution to the daily log-returns
3. Compares stable vs normal PDF graphically (log-scale)
4. Computes VaR at 1% and 5%
5. Performs a backtest: how many days exceeded the estimated VaR?


Suggested prompt

□ *"Write Python code using `scipy.stats.levy_stable` to fit a stable distribution to BTC-USD daily log-returns from `yfinance`. Plot histogram with stable and normal overlay on log-scale. Calculate VaR at 1% and 5% for both distributions and backtest against actual returns."*

Histogram + Stable vs Normal Fit


SFM_ch2_stable_fit

S&P 500 Log-Returns: Stable vs Normal Fit



PDF Comparison on Logarithmic Scale

SFM_ch2_stable_fit



Observations:

- ▶ Panel B (log scale) highlights tail differences
- ▶ $\alpha = 0.5$ (Lévy) has the heaviest tails
- ▶ At $\alpha = 2.0$ we recover the normal distribution exactly
- ▶ Financial returns typically have $\alpha \in [1.5, 1.9]$

VaR: Comparison Table


VaR on S&P 500 (2000–2025): Gaussian vs Stable vs Empirical

VaR level	Gaussian	Stable	Empirical
5%	−1.91%	−2.14%	−2.08%
1%	−2.71%	−4.12%	−3.56%
0.1%	−3.59%	−8.97%	−6.88%

- At the 5% level, differences are small
- At 1%, the stable model is $\sim 50\%$ more conservative than Gaussian
- At 0.1%, the stable model estimates losses **2.5** $\times$ larger
- Empirical VaR falls **between** the two parametric models

Cross-Asset Comparison: Stability Index α

SFM_ch2_stable_fit



Bitcoin has the lowest $\alpha \Rightarrow$ heaviest tails

Bonds have the highest α (closest to Gaussian)

Homework: Stable Distributions on BVB

Requirements

1. Choose a stock listed on the Bucharest Stock Exchange (e.g. TLV, SNP, BRD, FP)
2. Download daily data for the last 5 years
3. Fit a stable distribution to the log-returns
4. Compare stable vs normal PDF graphically (log scale)
5. Compute VaR at 1% and 5% (stable vs normal vs empirical)
6. Interpret the results: is the Romanian market “more Gaussian” than S&P 500?

Hints

- ▣ Use `yfinance` with the BVB ticker (e.g. `TLV.R0`)
- ▣ `levy_stable.fit()` can be slow — use `floc=np.mean(data)`
- ▣ Compare the estimated α with that of S&P 500

Competence Checklist

Theory

- ☐ Stability definition via sums
- ☐ Parameters: $\alpha, \beta, \gamma, \delta$
- ☐ Special cases: Gauss, Cauchy, Lévy
- ☐ Characteristic function
- ☐ Moment existence ($p < \alpha$)
- ☐ Tail behaviour (power law)

Practice

- ☐ McCulloch estimation (quantiles)
- ☐ MLE with `levy_stable.fit()`
- ☐ Koutrouvelis regression (ECF)
- ☐ VaR: stable vs normal vs empirical
- ☐ Simulation and visualisation in Python
- ☐ Interpreting α on real data

Seminar Planning

Week	Topic	Associated lecture	Main content
1	Returns	Ch. 1: Data sources	Return types, stationarity
2	Stable distributions	Ch. 2: Distributions	Stable fit, VaR, heavy tails
3	EMH & Random Walk	Ch. 3–4	EMH tests, variance ratio
4	Stylized facts	Ch. 5–6	ACF, volatility clustering

Quiz: 5 Questions

1. What are the bounds of the stability parameter α ?
 - (a) $0 < \alpha \leq 1$
 - (b) $0 < \alpha \leq 2$
 - (c) $-2 \leq \alpha \leq 2$
 - (d) $\alpha > 0$
2. For a distribution $S(1.6, 0, 1, 0)$, which moments **exist**?
 - (a) Mean and variance
 - (b) Only the mean
 - (c) All moments
 - (d) No moments
3. Which distribution is obtained at $\alpha = 1, \beta = 0$?
 - (a) Gaussian
 - (b) Cauchy
 - (c) Lévy
 - (d) Exponential

Quiz: 5 Questions (cont.) + Answers

4. In McCulloch's method, $v_\alpha = 2.44$ corresponds to:
- (a) Cauchy
 - (b) Gaussian
 - (c) Lévy
 - (d) Stable($\alpha = 1.5$)
5. Why is stable VaR > Gaussian VaR at 0.1%?
- (a) The stable distribution has heavier tails
 - (b) The stable distribution is more symmetric
 - (c) The Gaussian has infinite variance
 - (d) The parameter β is negative

Answers

- 1. (b) $0 < \alpha \leq 2$
- 2. (b) Only the mean (since $p < \alpha = 1.6$, only $p < 1.6$)
- 3. (b) Cauchy
- 4. (b) Gaussian ($\alpha = 2.0$)
- 5. (a) Heavier tails of the stable distribution

References

Textbooks

- ▣ J.P. Nolan, *Stable Distributions: Models for Heavy Tailed Data*, Birkhäuser, 2020.
- ▣ B. Mandelbrot, "The variation of certain speculative prices," *Journal of Business*, 36(4), 1963.
- ▣ S.T. Rachev, S. Mitnik, *Stable Paretian Models in Finance*, Wiley, 2000.
- ▣ J. Franke, W.K. Härdle, C.M. Hafner, *Statistics of Financial Markets*, 4th ed., Springer, 2019.
- ▣ J.H. McCulloch, "Simple consistent estimators of stable distribution parameters," *Communications in Statistics — Simulation and Computation*, 15(4), 1986.

Online resources

- ▣ <https://quantlet.com> — Quantlet repository
- ▣ <https://danpele.github.io/SFM/> — Course materials

Thank you! Questions? Seminar materials available at:

<https://danpele.github.io/SFM/>  Quantlet  Quantinar